

GEMM + All-Gather: Roofline vs PyTorch vs LLM CUDA Solution

